

# Assignment 03

Ava Godsy

September 22, 2025

## 1 Load the dataset

```
import pandas as pd
import plotly.express as px
import plotly.io as pio
from pyspark.sql import SparkSession
import re
import numpy as np
import plotly.graph_objects as go
from pyspark.sql.functions import col, split, explode, regexp_replace, transform, when
from pyspark.sql import functions as F
from pyspark.sql.functions import col, monotonically_increasing_id

np.random.seed(42)

pio.renderers.default = "notebook"

# Initialize Spark Session
spark = SparkSession.builder.appName("LightcastData").getOrCreate()

# Load Data
df = spark.read.option("header", "true").option("inferSchema", "true").option("multiLine", "true").text("lightcast_data.txt")
df.createOrReplaceTempView("job_postings")

# Show Schema and Sample Data
# print("----This is Diagnostic check, No need to print it in the final doc----")

# df.printSchema() # comment this line when rendering the submission
# df.show(5)
```

WARNING: Using incubator modules: jdk.incubator.vector

Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

25/09/24 03:46:52 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform  
java classes where applicable

25/09/24 03:47:07 WARN SparkStringUtils: Truncated the string representation of a plan since

## 2 Data Cleaning

```
df = df.withColumn("SALARY_FROM", col("SALARY_FROM").cast("float")) \
      .withColumn("SALARY_TO", col("SALARY_TO").cast("float")) \
      .withColumn("SALARY", col("SALARY").cast("float")) \
      .withColumn("MAX_YEARS_EXPERIENCE", col("MAX_YEARS_EXPERIENCE").cast("float")) \
      .withColumn("MIN_YEARS_EXPERIENCE", col("MIN_YEARS_EXPERIENCE").cast("float"))

def compute_median(sdf, col_name):
    q = sdf.approxQuantile(col_name, [0.5], 0.01)
    return q[0] if q else None

median_from = compute_median(df, "SALARY_FROM")
median_to = compute_median(df, "SALARY_TO")
median_salary = compute_median(df, "SALARY")

print("Medians:", median_from, "-", median_to, "-", median_salary)
```

[Stage 4:>

(0 + 1) / 1]

```
df = df.fillna({
    "SALARY_FROM": median_from,
    "SALARY_TO": median_to
})

df = df.withColumn("Average_Salary", (col("SALARY_FROM") + col("SALARY_TO")) / 2)
df = df.withColumn("EDUCATION_LEVELS_NAME", regexp_replace("EDUCATION_LEVELS_NAME", "\n ", " "))
df = df.withColumn("EDUCATION_LEVELS_NAME", regexp_replace("EDUCATION_LEVELS_NAME", "\n", " "))
df = df.withColumn("EDUCATION_LEVELS_NAME", regexp_replace("EDUCATION_LEVELS_NAME", "\r ", " "))
df = df.withColumn("EDUCATION_LEVELS_NAME", regexp_replace("EDUCATION_LEVELS_NAME", "\r", " "))
```

```

export_cols = [
    "EDUCATION_LEVELS_NAME",
    "REMOTE_TYPE_NAME",
    "MAX_YEARS_EXPERIENCE",
    "Average_Salary",
    "LOT_V6_SPECIALIZED_OCCUPATION_NAME",
    "NAICS2_NAME",
    "EMPLOYMENT_TYPE_NAME",
    "ONET_NAME",
    "SALARY_FROM",
    "SALARY_TO",
    "SALARY"
]
df_selected = df.select (*export_cols)

pdf = df_selected.toPandas()
pdf.to_csv("lightcast_cleaned.csv", index=False)

print(" Data cleaning complete. Rows retained:", len(pdf))
# df_selected.show(5)

```

Data cleaning complete. Rows retained: 72498

### 3 Salary Distribution by Industry and Employment Type

The following data

```

import plotly.express as px

fig = px.box(
    df_selected,
    x="NAICS2_NAME",
    y="SALARY_FROM",
    color="EMPLOYMENT_TYPE_NAME",
    title="Salary Distribution by Industry and Employment Type",
    points="all", # Show all points
    notched=True, # Notched boxes
    height=1000, # Taller figure

```

```

    color_discrete_sequence=["purple", "blue", "green"] # Custom colors
)

fig.update_layout(
    title_font=dict(family="Garamond", size=24, color="black"),
    xaxis_title="Industry (NAICS2)",
    yaxis_title="Starting Salary",
    boxmode="group", # Grouped box plots
    xaxis_tickangle=45, # Rotate x-axis labels
    font=dict(
        family="Garamond, serif", # Set font to Garamond
        size=12
    )
)

fig.show()

```

Unable to display output for mime type(s): text/html

## 4 Salary Analysis by ONET Occupation Type (Bubble Chart)

```

from pyspark.sql import functions as F

df_filtered = df_selected.filter(F.col("SALARY").isNotNull())

lot_salary = df_filtered.groupBy("LOT_V6_SPECIALIZED_OCCUPATION_NAME").agg(
    F.expr("percentile_approx(SALARY, 0.5)").alias("Median Salary"),
    F.count("*").alias("Job_Postings")
)

lot_salary.show()

```

[Stage 8:>

(0 + 1) / 1]

LOT_V6_SPECIALIZED_OCCUPATION_NAME	Median Salary	Job Postings
Business Intellig...	107500.0	1800
Business Analyst ...	93650.0	1640
Healthcare Analyst	89440.0	94
Oracle Consultant...	138750.0	3526
SAP Analyst / Admin	120640.0	3373
Data Analyst	96100.0	12377
General ERP Analy...	125900.0	3703
Marketing Analyst	94500.0	65
Enterprise Architect	157600.0	3321
Financial Data An...	49920.0	429
Data Quality Analyst	96600.0	480

```
import plotly.express as px

fig = px.scatter(
    lot_salary,
    x="LOT_V6_SPECIALIZED_OCCUPATION_NAME",
    y="Median Salary",
    size="Job Postings",
    color="Median Salary",
    hover_name="LOT_V6_SPECIALIZED_OCCUPATION_NAME",
    size_max=60,
    title="Salary Analysis by LOT Occupation Type",
)
fig.update_layout(
    title_font=dict(family="Garamond", size=24, color="black"),
    font=dict(family="Garamond", size=12, color="black"),
    plot_bgcolor="white",
    paper_bgcolor="#f7f7f7",
    xaxis=dict(title="Occupation Name", tickangle=45),
    yaxis=dict(title="Median Salary ($)", gridcolor="#e5e5e5"),
)

# Show the figure
fig.show()
```

Unable to display output for mime type(s): text/html

## 5 Salary by Education Level

```
# df_selected.select("EDUCATION_LEVELS_NAME").distinct().show(truncate=False)

from pyspark.sql.functions import col, when

# Create the EDU_GROUP column based on EDUCATION_LEVELS_NAME
df_with_edu_group = df_selected.withColumn(
    "EDU_GROUP",
    when(
        col("EDUCATION_LEVELS_NAME").rlike("(?i)No Education Listed|GED|Associate"),
        "Associate's or lower"
    ).when(
        col("EDUCATION_LEVELS_NAME").rlike("(?i)Bachelor"),
        "Bachelor's"
    ).when(
        col("EDUCATION_LEVELS_NAME").rlike("(?i)Master"),
        "Master's"
    ).when(
        col("EDUCATION_LEVELS_NAME").rlike("(?i)Ph\\.D\\.|professional degree"),
        "PhD"
    ).otherwise("Associate's or lower") # Optional: handle unmatched entries
)

# Select required columns
final_df = df_with_edu_group.select(
    "EDU_GROUP",
    "LOT_V6_SPECIALIZED_OCCUPATION_NAME",
    "Average_Salary",
    "MAX_YEARS_EXPERIENCE"
)

final_df.show()
```

```
+-----+-----+-----+-----+
+-----+
|          EDU_GROUP|LOT_V6_SPECIALIZED_OCCUPATION_NAME|Average_Salary|MAX_YEARS_EXPERIENCE
```

+-----+-----+-----+			
+-----+-----+-----+			
Bachelor's	General ERP Analy...	108668.5	2.0
Associate's or lower	Oracle Consultant...	108668.5	3.0
Bachelor's	Data Analyst	108668.5	NULL
Associate's or lower	Data Analyst	108668.5	NULL
Associate's or lower	Oracle Consultant...	92500.0	NULL
Bachelor's	Data Analyst	110155.0	NULL
Bachelor's	Data Analyst	108668.5	NULL
Bachelor's	Data Analyst	108668.5	NULL
Associate's or lower	General ERP Analy...	108668.5	7.0
Bachelor's	Data Analyst	92962.0	2.0
Associate's or lower	Data Analyst	107645.5	NULL
Associate's or lower	Data Analyst	108668.5	NULL
Bachelor's	Data Analyst	108668.5	NULL
Bachelor's	General ERP Analy...	192800.0	NULL
Associate's or lower	Enterprise Architect	81286.0	NULL
Associate's or lower	Data Analyst	108668.5	5.0
Associate's or lower	General ERP Analy...	125900.0	NULL
Associate's or lower	Oracle Consultant...	108668.5	3.0
Bachelor's	Enterprise Architect	165000.0	8.0
Associate's or lower	Data Analyst	170000.0	NULL
+-----+-----+-----+			
+-----+-----+-----+			

only showing top 20 rows

```
import plotly.express as px
import numpy as np

# Step 1: Convert PySpark DataFrame to Pandas
pdf = final_df.toPandas()

# Step 2: Add jitter to MAX_YEARS_EXPERIENCE
np.random.seed(0)
jitter_strength = 0.1
pdf['JITTERED_EXPERIENCE'] = pdf['MAX_YEARS_EXPERIENCE'] + np.random.uniform(
    -jitter_strength, jitter_strength, size=len(pdf)
)

# Step 3: Define custom color mapping
color_map = {
    "Associate's or lower": 'yellow',
```

```

    "Bachelor's": 'green',
    "Master's": 'blue',
    "PhD": 'purple'
}

# Step 4: Create the Plotly scatter plot
fig = px.scatter(
    pdf,
    x='JITTERED_EXPERIENCE',
    y='Average_Salary',
    color='EDU_GROUP',
    color_discrete_map=color_map,
    title="Salary by Education Level",
    labels={
        'JITTERED_EXPERIENCE': 'Max Years of Experience Required (jittered)',
        'Average_Salary': 'Average Salary',
        'EDU_GROUP': 'Minimum Education Level Required'
    },
    opacity=0.7
)

# Step 5: Update layout with Garamond font and sizes
fig.update_layout(
    title_font=dict(family='Garamond', size=24, color='black'),
    font=dict(family='Garamond', size=12, color='black'),
    legend_title_font=dict(family='Garamond', size=12, color='black'),
    legend_font=dict(family='Garamond', size=12, color='black')
)

# Step 6: Show the figure
fig.show()

```

Unable to display output for mime type(s): text/html



## 6 Salary by Remote Work Type

```
# df_selected.select("REMOTE_TYPE_NAME").distinct().show(truncate=False)

from pyspark.sql.functions import col, when

df_with_remote_group = df_selected.withColumn(
    "REMOTE_GROUP",
    when(
        col("REMOTE_TYPE_NAME") == "Remote", "Remote"
    ).when(
        col("REMOTE_TYPE_NAME") == "Hybrid Remote", "Hybrid"
    ).when(
        (col("REMOTE_TYPE_NAME").isNull()) |
        (col("REMOTE_TYPE_NAME") == "Not Remote") |
        (col("REMOTE_TYPE_NAME") == "[None]"),
        "Onsite"
    ).otherwise("Onsite"))

remote_df = df_with_remote_group.select(
    "REMOTE_GROUP",
    "LOT_V6_SPECIALIZED_OCCUPATION_NAME",
    "Average_Salary",
    "MAX_YEARS_EXPERIENCE"
)

remote_df.show()
```

```
+-----+-----+-----+-----+
+
|REMOTE_GROUP|LOT_V6_SPECIALIZED_OCCUPATION_NAME|Average_Salary|MAX_YEARS_EXPERIENCE|
+-----+-----+-----+-----+
+
|    Onsite|General ERP Analy...|    108668.5|          2.0|
|    Remote|Oracle Consultant...|    108668.5|          3.0|
|    Onsite|          Data Analyst|    108668.5|        NULL|
|    Onsite|          Data Analyst|    108668.5|        NULL|
|    Onsite|Oracle Consultant...|     92500.0|        NULL|
|    Remote|          Data Analyst|   110155.0|        NULL|
|    Onsite|          Data Analyst|    108668.5|        NULL|
|    Onsite|          Data Analyst|    108668.5|        NULL|
```



```

    title="Salary by Remote Status",
    labels={
        'JITTERED_EXPERIENCE': 'Max Years of Experience Required',
        'Average_Salary': 'Average Salary',
        'REMOTE_GROUP': 'Remote Status'
    },
    opacity=0.7
)

# Step 5: Update layout with Garamond font and sizes
fig.update_layout(
    title_font=dict(family='Garamond', size=24, color='black'),
    font=dict(family='Garamond', size=12, color='black'),
    legend_title_font=dict(family='Garamond', size=12, color='black'),
    legend_font=dict(family='Garamond', size=12, color='black')
)

# Step 6: Show the figure
fig.show()

```

Unable to display output for mime type(s): text/html

```

# Step 1: Convert to Pandas
pdf = remote_df.toPandas()

# Step 2: Create plot using Plotly
import plotly.express as px

color_map = {
    "Remote": 'yellow',
    "Hybrid": 'green',
    "Onsite": 'blue',
}

fig = px.histogram(
    pdf,
    x="MAX_YEARS_EXPERIENCE",
    y="Average_Salary",
    color="REMOTE_GROUP",

```

```

color_discrete_map=color_map,
histfunc="avg",
nbins=int(pdf['MAX_YEARS_EXPERIENCE'].max()) + 1,
barmode='group',
title="Average Salary by Years of Experience and Remote Type",
labels={
    'MAX_YEARS_EXPERIENCE': 'Max Years of Experience Required',
    'Average_Salary': 'Average Salary',
    'REMOTE_GROUP': 'Remote Status'
}
)

fig.update_layout(
    title_font=dict(family='Garamond', size=24, color='black'),
    font=dict(family='Garamond', size=12, color='black'),
    legend_title_font=dict(family='Garamond', size=12, color='black'),
    legend_font=dict(family='Garamond', size=12, color='black'),
    xaxis=dict(dtick=1),
    yaxis_title="Average Salary",
    bargap=0.2
)

fig.show()

```

Unable to display output for mime type(s): text/html